PH1102E

Week 4

Duties to Strangers

I. Singer's argument

There is no doubt that, for the amount of money that it takes to buy a car in Singapore, you could save someone's life. For that amount of money, you could buy a suitcase full of malaria medication, personally transport it to a needy village in rural Africa (or Haiti, or wherever), and distribute it yourself. Better yet, you could donate the same amount of money to a well-managed relief organization, which would use it to provide even more people with life-saving medicines. Given that the medicine to cure a case of malaria only costs about $10, and given that a car in Singapore typically costs more than $100,000, even if 99.99% of your donation gets absorbed by administrative overhead, enough of it would still get through to save someone's life.

Singapore has an excellent system of public transport. Few people living here need to have a car in order to survive, get to work, etc. Yet there are hundreds of thousands of privately owned cars in Singapore. Each of these cars corresponds to at least one human life that could have been saved, but was not. Yet we do not think of buying a car as an immoral act, and we do not think of car owners as moral monsters. Now, any respectable relief organization converts a lot more than .01% of the funds it receives into aid for those in need (like rural African victims of malaria). Most of the biggest ones (Oxfam, Red Cross, etc.) convert over 75% of the funds they receive into actual aid (food, medicine, or whatever). $10 is 75% of $13.34. So, if you were to donate $13.34 to an efficient relief organization specializing in the treatment of diseases like malaria, enough of your donation would get through to the people who need it to save someone's life. Yet, if we don't think that it is immoral to buy a car, even less so do we think that it is immoral to spend $20 on a bottle of wine, or a DVD, or a birthday present for a friend.

Peter Singer thinks that our ordinary ways of thinking about these things are incorrect. He thinks that it is immoral to buy a car, a bottle of wine, or anything else that could be categorized as a "luxury." He argues for this seemingly radical view as follows.

Suppose you are on your way to the post office, to pay your monthly utilities bill. The post office closes in half an hour, but you have just enough time to get there if you hurry. If you don't pay the bill today, you will be assessed a late payment fee of $20. As you are hurrying along, you pass a pond in which a child is struggling. The pond is small -- you could easily swim across it without danger -- but it is obvious that the child will drown unless you stop to rescue it. Of course, if you stop to rescue the child, you'll have to pay the $20 penalty for late payment of your utilities bill.

Singer contends that it would be immoral -- in fact, morally outrageous -- for you not to stop and save the child, in this situation. He also contends that if you don't save the child, your behavior is in all relevant respects identical to what it would be if you spent $20 on a bottle of wine (or a CD, or whatever) , instead of donating it to a reputable relief organization. Singer's conclusion: it is immoral -- in fact, morally outrageous -- to purchase a CD or a bottle of wine, instead of contributing the money to some such organization.

If Singer is correct, each of us engages in morally outrageous behavior on a daily basis -- many of us more than once a day. Perhaps the $20 figure is too low? Well, suppose that it takes a contribution of

$1000 save someone's life. And now suppose that the consequence of saving the child in the pond will be that you have to pay a late fee of $1000, instead of $20. Most of us would still consider it immoral not to save the child, even though the cost of doing so now is significantly higher than in the original example. But if it would be immoral to let the child die despite the fact that it would cost you $1000 to save it, it is, Singer contends, equally immoral not to contribute $1000 to a suitable relief effort. Yet most of us spend at least that much every year on luxuries (ice cream, extra clothing, movies, etc.).

There are certain nuances here that we need to be aware of. Suppose that you are rushing to the post office to pay your bills on time so that you will not lose $1000 which you have saved up to donate to a relief organization that will use your donation to save multiple lives. In this case what should you do? Well, according to Singer, you should let the child drown. If you save the child, you'll forfeit the $1000 you had saved up for the relief organization, as a result of which the organization will not receive that $1000 donation, as a result of which more than one child will die (of malaria, or malnutrition, or some other preventable ill). But this rationale for letting the child drown depends on your actually donating the money to the relief organization. If you don't donate the money, your failure to save the child is still morally monstrous -- according to Singer.

Simply stated, Singer's argument is this:

(1) If you spend any significant amount of money on luxuries, you do something that is in all relevant respects equivalent to failing to rescue the drowning child.

(2) It would be profoundly wrong for you not to rescue the drowning child.

(3) Therefore, if you spend any significant amount of money on luxuries, you do something that is profoundly wrong.

Most of Singer's critics focus on the first premise of the argument (step 1), but it is also possible to object to its second premise. We'll consider both kinds of objections, beginning with objections to premise (1).

II. Objections to Premise (1)

Most people who wish to resist Singer's conclusion focus on premise (1). Here, the strategy is to concede to Singer that it would be morally wrong not to rescue the drowning child, but to argue (against Singer) that there is a morally relevant difference between failing to rescue the child, and failing to forego luxuries in order to support global relief efforts. Let's look at three of the more prominent objections along these lines.

A. Can't save them all

No matter how much of your wealth you donate to the effort to combat malaria in Africa, hundreds of children will die of malaria in Africa every day. The problem is simply too large for you to solve. By

contrast, if you jump into the pond, you solve the problem of the drowning child completely, just by saving his life.

This is a weak objection. Suppose that the pond contained many children, not just one. All of them are drowning, but you can't save them all: there are too many of them, and the pond is too large. Does this mean that there would be nothing wrong with your failing to jump in? Surely not. From the fact that you cannot save all the children in pond, it doesn't follow that you have no obligation to save any of them. No one is going to blame you for failing to save all the children -- that, we are supposing, was humanly impossible -- but we would still blame you if you failed to save any of them (as many as you could, given your swimming abilities, the size of the pond, etc.).

A related objection to premise (1) is that since so many African children are suffering from malaria, there is no particular child who will die if I make no donation, whereas there is a very particular child who will die if I do not jump into the pond.

Now, it is true that in the pond case, I can point to a specific child and say: "That is the child I can save," whereas there is no African child to whom I can point, saying the same thing. But from this it does not follow that there is not a very specific African child who will die, if I do not make a donation. It's not just that some-child-or-other will die of malaria, or that some "unspecific" or "indeterminate" child will die (whatever that means). When I fail to donate, that much less money enters the relief organization's pipeline, and that much less comes out the other end, in the form of medicine. Suppose I do donate. In theory, we could track my contribution through the organization's inventory and accounting system, and identify exactly which pills would not have been delivered when needed, and exactly which African children would not have received pills when they needed them, had I not made my donation. By the same token, if I do not donate, and if there is an all-knowing God, he could take me to a cemetery in an African village and point to the precise graves of the children who died of malaria, but wouldn't have died, had I donated. So, there is no genuine disanalogy here between my relationship to the child in the pond, and my relationship to the child in Africa. It is just that the relationship is more open to view in the former case than in the latter.

B. My personal obligation is small

A somewhat more plausible-sounding objection to premise (1) is that when it comes to the Africans, there are millions of people like me who are equally in a position to help, whereas in the pond scenario, only I can rescue the child, there being no one else around. That being so, my personal obligation to contribute to famine relief etc. is small: it is but one sliver of a much larger obligation that all relatively wealthy people shoulder collectively.

But this objection too collapses under scrutiny. Suppose that I am walking to the post office with a whole crowd of people, all of whom need to pay their utility bills in order to avoid late fees. We all see the child in the pond. If we all walk on, letting the child drown, is my behavior any better than in the original scenario? It seems pretty clear that it is not.

Now, I might be forgiven for glancing around at the other people before I jump in, to see if anyone else is going to jump in first. There's no point in my getting wet if someone else is already on the job. But if it is obvious that no one else is going to jump in, then surely it would be wrong for me not to jump in myself. And this is the situation I face when it comes to the decision whether or not to contribute to global relief efforts. I know darn well that these organizations are not going to receive enough money to rescue everyone who needs rescuing (from famine, disease, or what have you). The millions of other first-world citizens who are just like me in terms of their ability to donate are analogous to the people standing around the pond who I can see are clearly not going to jump in. So, once again, we find no important disanalogy between my relationship to the child in the pond, and my relationship to the children in Africa.

C. Distance matters

The most obvious difference between the child-in-the-pond case and the child-in-Africa case is that in the former, but not the latter, the physical distance between me and the person who stands in need of assistance is small. The African children are far away; the child in the pond is right in front of me. The question is whether this is a relevant difference: a difference that is relevant to the moral status of my behavior, in each case.

One might have thought that the fact that I'm physically close to the child in the pond makes it easier for me to rescue him, and correspondingly worse of me not to make the (relatively small) effort required to do so. But in this day and age, it may actually require more effort to pull a child out of a pond than to provide a far-away child with life-saving medicine. To accomplish the former, I have to exert myself physically, and get wet. (We could even suppose that I have to scale a chain-link fence to get to the pond.) To accomplish the latter, I only have to point my web browser to the website of an appropriate relief agency, and key in my credit card particulars. I don't even have to get out of my chair.

So if distance makes a moral difference, it is not because farther away people are more difficult to save than nearby ones. But how else could distance make a morally relevant difference?

Some philosophers argue that it makes a difference for the following reason. If I neglect to rescue the child, I prove myself to have a deeply flawed, even repulsive, character. I prove that I am the kind of person who can look at a drowning child, knowing that I could easily rescue it, and walk on, just to save myself a little bit of money. By contrast, if I neglect to donate to an organization that fights deadly childhood diseases in far-flung parts of the globe, I do not demonstrate such a flawed character. My failure to donate proves that I have an imperfect character, but not that my character is repulsive. If we think of the wrongness of an action as a matter of how flawed the character is from which it flows, we can say that my failure to save the drowning child is a far more serious instance of wrongdoing than my failure to give to charity, since it proceeds from a more seriously flawed character.

The idea that the moral worth (positive or negative) of an action depends on the quality of the character from which it flows has some plausibility to it. It is the main alternative to utilitarianism: whereas utilitarianism focuses on the effects of actions (in terms of how much happiness or suffering they cause), the "character theory" of morality focuses on the actions' causes (the motives and intentions of the

people who perform the actions). The problem, however, is that it is hard to see how a failure to donate to charity, knowing the consequences of one's failure to donate, is evidence of a character any less flawed than that of someone who lets a child drown, as in Singer's example.

Indeed, one might even argue that the man who lets the child drown proves himself to have a character that is in some respects superior to that of the man who refuses to donate to relief efforts. At least the man who passes by the pond faces up to the consequences of his actions directly, whereas the man who refuses to donate most likely does his best to put out of his mind the deaths that he allows to occur. And anyway, if my action of rescuing the child from the pond really proceeds from a good character, how can it be that the same character fails to prompt me to donate to charity (or "charity"), knowing full well that the consequences of failure to do so are no less dire than those of failing to save the drowning child? If I take action in one case but not the other, it looks like I am motivated more by squeamishness than by an abiding concern for human welfare.

III. Objections to Premise (2)

So far we have considered various attempts to undercut the first premise of Singer's argument, by identifying some relevant disanalogy between the case of the drowning child, and the case of the child stricken with malaria. None of these has been successful. Let's consider some attempts to cast doubt on Singer's second premise.

A. No duty to rescue

One objection to premise (2) is that you have no duty to rescue the child. True, it would be nice of you to save it, but the child has no right to be saved by you: it's not as if you owe it to the child to save it, as you would if the child were in the pond because you threw it in there. But a duty implies a right: I have a duty to rescue the child only if the child has a right to be rescued by me. Since the child has no such right, I have no such duty.

The problem with this objection is that even if the child has no right to my intervention, and even if it follows that I have no duty to rescue the child, still, the fact remains that I would be behaving very badly if I didn't rescue it. Whether we want to characterize this badness in terms of a violation of duty or in some other way, the basic fact remains: if I don't rescue the child, I do something wrong. If I don't violate any of my duties, that just proves that I can do something very wrong without violating any of my duties.

Let's consider a different objection to premise (2).

B. Cost and value

How much is a human life worth? One possible answer is that a human life has infinite worth, infinite value. But this cannot be. If a human life had infinite value, then it would make sense to sacrifice a million lives to save one (a million times infinity = infinity).

Given that each human life has only some finite value, it only makes sense to sacrifice a certain amount of value in order to save that life. The question is: how do we determine the value of a life?

One way is in economic terms. At least, we can do this as long as we are talking about utilitarian value (which is the only kind of value that Singer himself recognizes). The utilitarian value of a person's life is basically the amount of net happiness that that person creates over the course of his lifetime. This value is at least approximated by the person's lifetime income. Your lifetime income is the amount of money people give you for what you give them. Ultimately, what people want is happiness. So, in effect, each of us gets paid to make other people happy. The connection with happiness is not always direct or obvious, but it's there. An entertainer gets paid to entertain his audience -- in this case, the connection with happiness is direct and obvious. Someone working in a pencil factory gets paid to produce pencils, which people then use to achieve certain things they want to achieve, and thereby make themselves (slightly) happier. The prison warden gets paid to help keep dangerous people off the streets, and thus indirectly to decrease the amount of suffering that society at large must endure.

(Of course, lifetime income can only be a very rough measure of a life's utilitarian value. A successful pickpocket might amass a vast fortune, but do far more harm than good in the process. Flawed as it is, however, lifetime income, or some similar measure of personal economic output, seems to be the best way we currently have of quantifying the utilitarian value of a life. And quantify it we must, if we are to implement the utilitarian code of conduct -- to do always what maximizes happiness, or minimizes suffering -- in an intelligent way.)

So, taking lifetime income as a rough approximation of the utilitarian value of a human life -- and remember, Singer himself is a devout utilitarian -- we can now ask: what is the prospective lifetime income of the child in Africa? The answer turns out to be: about $40,000. If the child survives and works for a normal span of years, that is how much it can expect to earn, in total, over the course of its entire life.

So the life of the African child is worth about $40,000, by this rough utilitarian reckoning. Does that mean that I should now be willing to sacrifice up to $40,000 in order to save the child? No. The $40,000 of lifetime income -- the $40,000 of utilitarian value -- that the child accumulates comes in over a period of 60 years (assuming that the child stops earning at around the age of 60). If I were to invest $40,000 today, then I could expect to have well over $2,000,000 at the end of 60 years. Investing the money would therefore create over 50 times as much utilitarian value as using it to save the African child. (The returns on my investment are income, which I receive by directly or indirectly making people happy -- people who buy the goods or services of the enterprises that I help to capitalize.)

But, realistically, a donation of far less than $40,000 would suffice to save the African child. The question is, what is the *most* that I can spend to save this child's life, without forgoing an opportunity to create more than $40,000 worth of (utilitarian) value?

Well, assuming that I can get an average 7% annual return on investment over a sixty year period -- which should not be hard to achieve at all -- I shouldn't be willing to spend more than about $700 to save the child. ($690.30 invested for sixty years at 7% annual return = $40,000.42.)

By the same token, I shouldn't be willing to sacrifice more than $700 to save the child in the pond. If I do, then I am not getting an optimal return on my investment, by utilitarian standards. Of course, given that the drowning child is in Singapore, its prospective lifetime income is likely much higher than the African child's -- say, something on the order of $50,000 *per year*, or about two million dollars over the course of a lifetime. Thus I should be willing to sacrifice up to about $40,000 in order to save the drowning child. This is just because the child's life is worth more, in utilitarian terms.

Anyway, the fact remains that I should be willing to spend up to $700 to save an African child's life. And there are thousands upon thousands of such lives in need of saving. So Singer's logic still seems to prevail: assuming that I have good reason to suppose that the amount of humanitarian aid people need this year (to stay alive) is greater than the amount of aid they'll actually receive, I have good reason to think that every dollar I donate will go toward saving a life that would otherwise end prematurely. Assuming furthermore that $700 is enough money to save a life -- remember that a dose of malaria medication only costs about $10 -- I ought to be prepared to spend many multiples of $700 to save these lives before I spend a penny on anything else.

These assumptions might be questioned. Perhaps there are hidden costs to providing humanitarian aid that push the average cost of saving an individual life above $700. Perhaps there is reason for optimism that the annual $7 billion dollar shortfall in aid (valuing each of the 10 million lives lost annually to malaria etc. at $700 each) will be closed by a modest increase in worldwide charitable giving. But if the cost of saving a life is less than the cost of failing to save it, and if historic shortfalls in aid continue into the future, Singer's argument still compels us to give far more to needy strangers than most of us currently do.

IV. Concluding remarks

There is no easy objection to Singer's argument, no obvious flaw in his reasoning. His basic, underlying thought is that we must save a life whenever we can do so without sacrificing anything of comparable value. Of the objections to Singer that we have considered above, probably the most promising is the objection that the lives we can save are of sufficiently little value that Singer's principle does not require us to devote more than a modest amount of our personal resources to famine relief, anti-malaria campaigns, etc. Still, given that most of us devote *none* of our resources to such efforts, even this might leave Singer's argument with enough force to compel us, morally, to change the way we live our lives in significant ways.